Openshift Installation On IBM Cloud Infrastructure

This document will help you to install Openshift on a VMWare based Infrastructure, It provides detailed steps you need to follow to access IBM Cloud, to provision the needed Infrastructure, and to install Openshift 3.11 on top of it.

This document is divided into three main sections:

1- How to access IBM Cloud using VPN.
2- How to provision ESXi Host. and how to use it to create the required VMs.
3- How to install Openshift 3.11
    3.1- Preparing the hosts (CentOS7 VMs will be used here)
    3.2- Installing Openshift using Ansible Playbook
    3.3- Troubleshooting

Section 01- IBM Cloud Access

When you provision an ESXi host on top of IBM Cloud, In order to access it, you will need to setup a VPN.
You will need to establish a VPN connection to IBM Cloud private network via SSL, or IPSec.

Pre-Req:

Identify which Endpoint you will use, Which DC you will provision your BM on. Here is a full list of end points
The environment I will use in here will be based on Frankfurt , which have endpoint vpn.fra02.softlayer.com

Based on your local machine, install your VPN Client
For Windows Users
For Mac Users

From a Mac machine, install MotionPro Plus using App Store

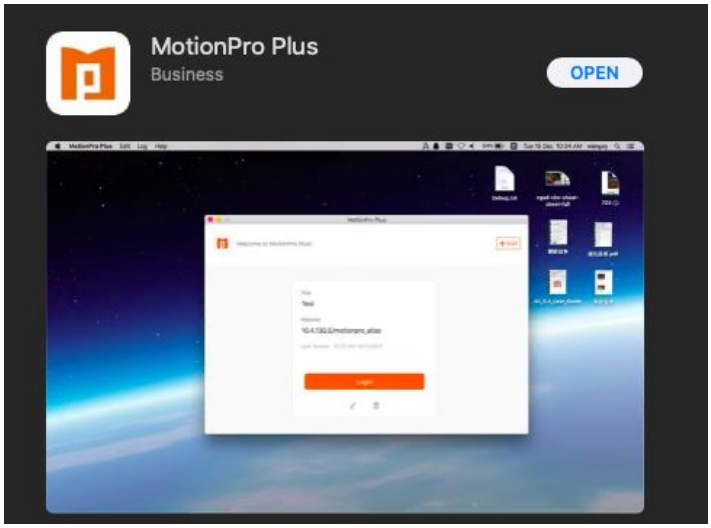

On IBM Cloud Portal, Enable SSL VPN Access for the user and manage its password

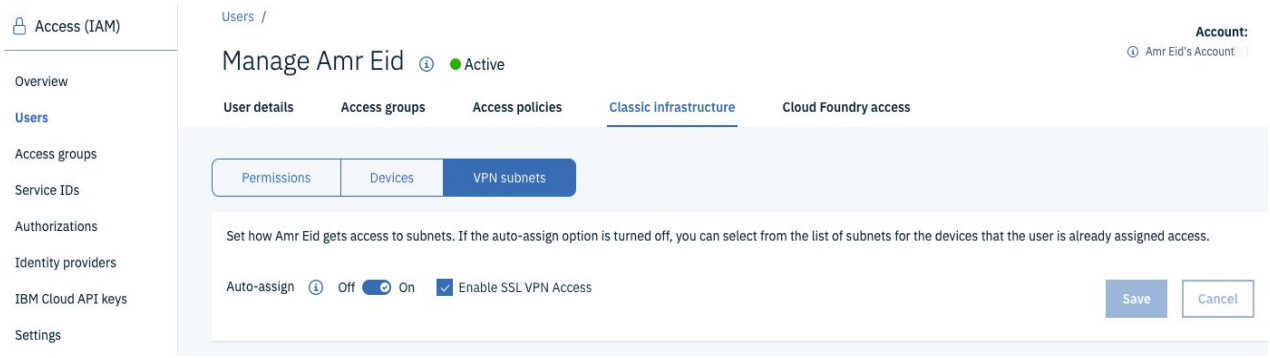

Preferably change the password. If not you need to capture the user/password from IAM/Users/User Page

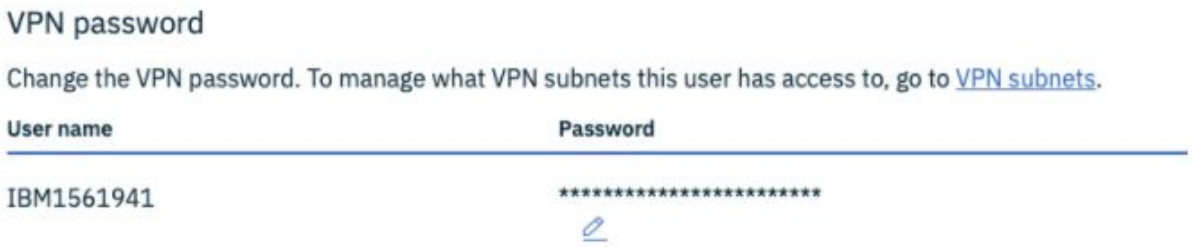

Get the Client up and running
You will need to feed MotionPro with the endpoint, In our case "vpn.fra02.softlayer.com"
You will also need to provide it with the user/password after you hit login.

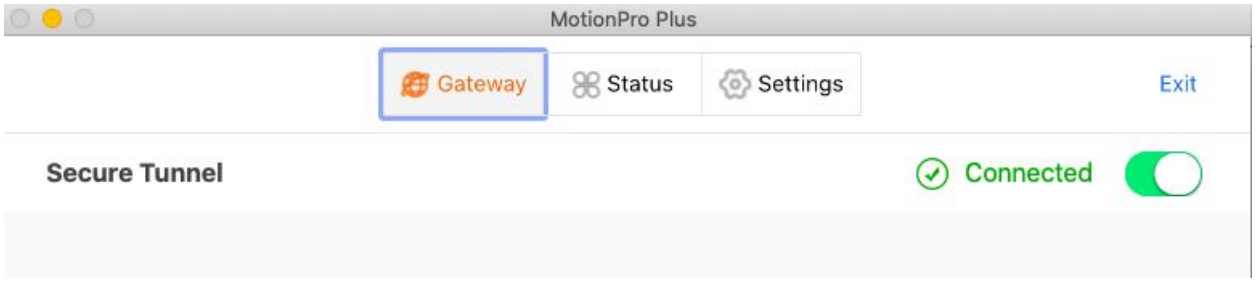

Section 02- ESXi Host Provisioning

Provisioning the Bare Metal usually takes from 8 to 24 hours, so plan ahead.
Form the portal, Pick the vmWare BM version you like.

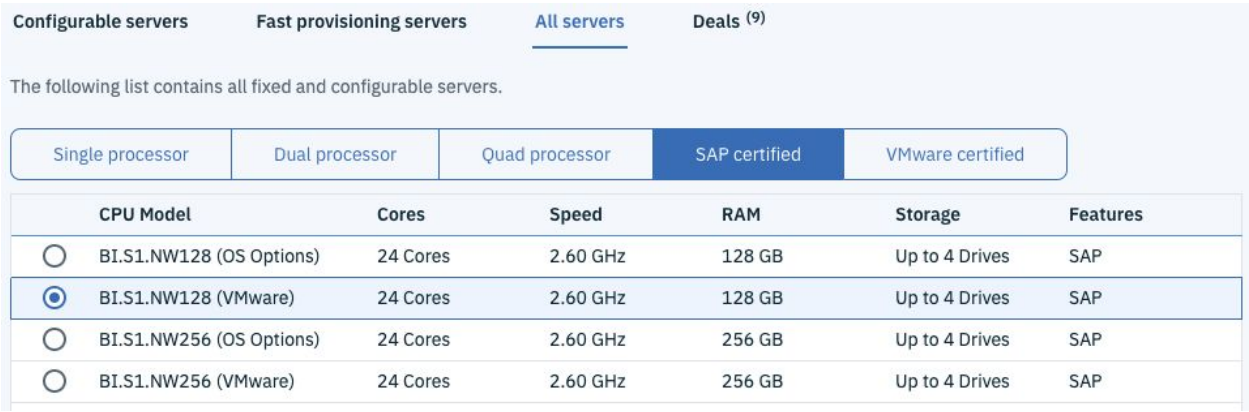

	CPU Model	Cores	Speed	RAM	Storage	Features
○	BI.S1.NW128 (OS Options)	24 Cores	2.60 GHz	128 GB	Up to 4 Drives	SAP
◉	BI.S1.NW128 (VMware)	24 Cores	2.60 GHz	128 GB	Up to 4 Drives	SAP
○	BI.S1.NW256 (OS Options)	24 Cores	2.60 GHz	256 GB	Up to 4 Drives	SAP
○	BI.S1.NW256 (VMware)	24 Cores	2.60 GHz	256 GB	Up to 4 Drives	SAP

Check the Subnets / IP Addresses attached to the BM, Pick the private IP (usually attached to eth0), the other private IP usually is used for the management and could be accessed through KVM.

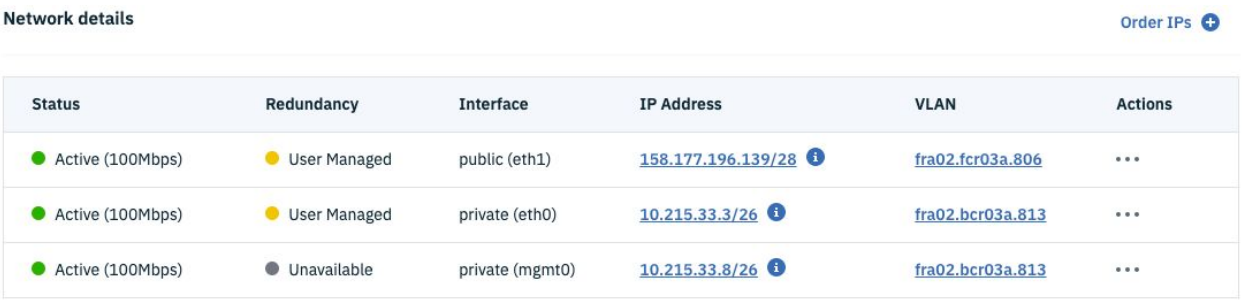

Status	Redundancy	Interface	IP Address	VLAN	Actions
Active (100Mbps)	User Managed	public (eth1)	158.177.196.139/28	fra02.fcr03a.806	•••
Active (100Mbps)	User Managed	private (eth0)	10.215.33.3/26	fra02.bcr03a.813	•••
Active (100Mbps)	Unavailable	private (mgmt0)	10.215.33.8/26	fra02.bcr03a.813	•••

From any browser type the private IP and hit enter, then accept the default certificates you may face. You should receive the login page like below.

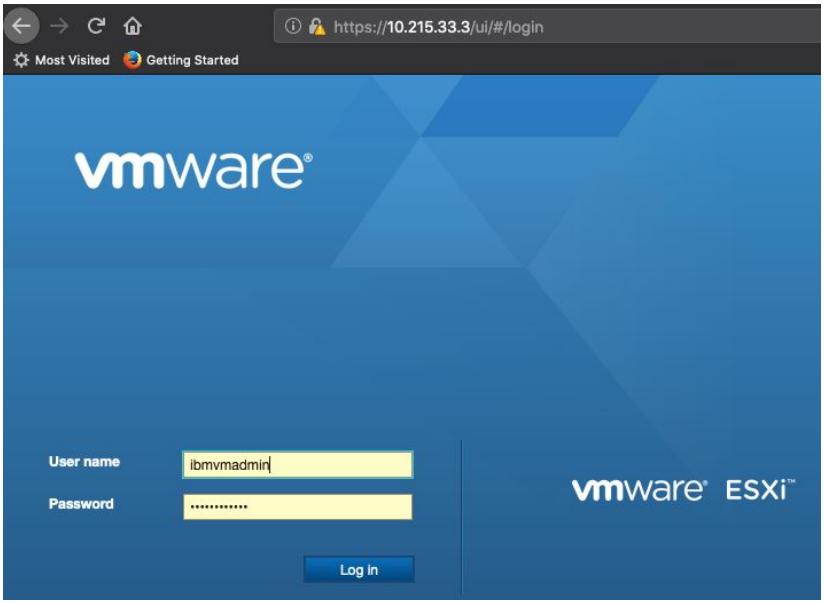

From your BM page on IBM Cloud portal, navigate to Passwords to get the vSphere username and password. (ibmvmadmin is the ESXi user).

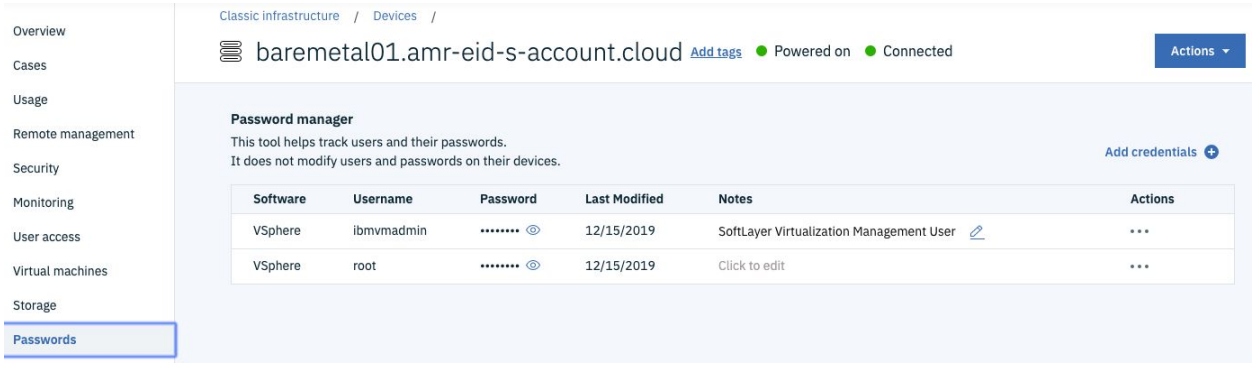

Configuring ESXi

In the ESXi portal , You need to know

- Physical NICs on the BM are reflected on the ESXi physical NICs view.
- You should have validate that you do have an up and running NICs like the below screenshot.

Here, There are two NICs down and the other two are up and running.
From the mapping you can easily detect that NIC0 is the private NIC and NIC1 is the Public.

In order to proceed , you just need to do the below actions

- Create Port Groups (one for the private and one for the public.)
- link each vSwitch with its related Physical NIC.

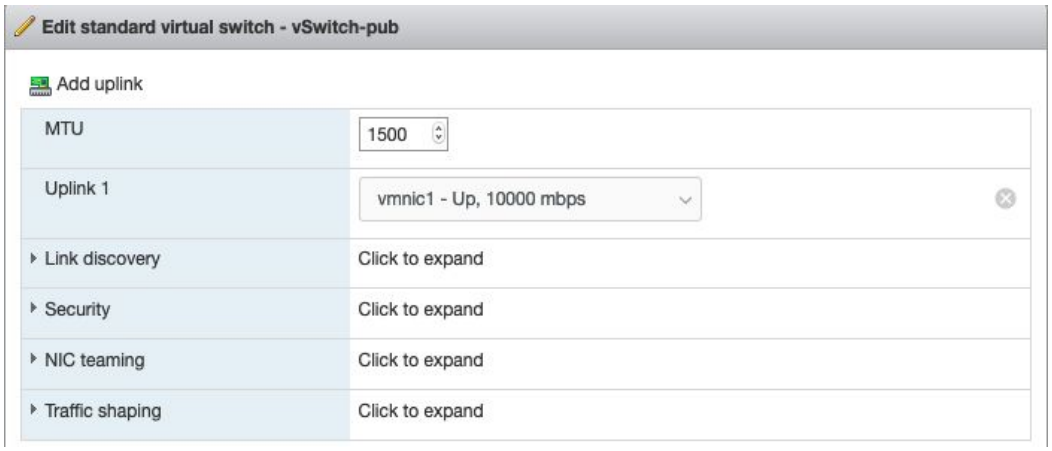

Below screenshot shows how the public looks, attached to vmnic1 which is the public interface.
Apply the same for the private and make sure it is connected to vmnic0
VMKernel NICs will have one for private connection, you will need to create one more.

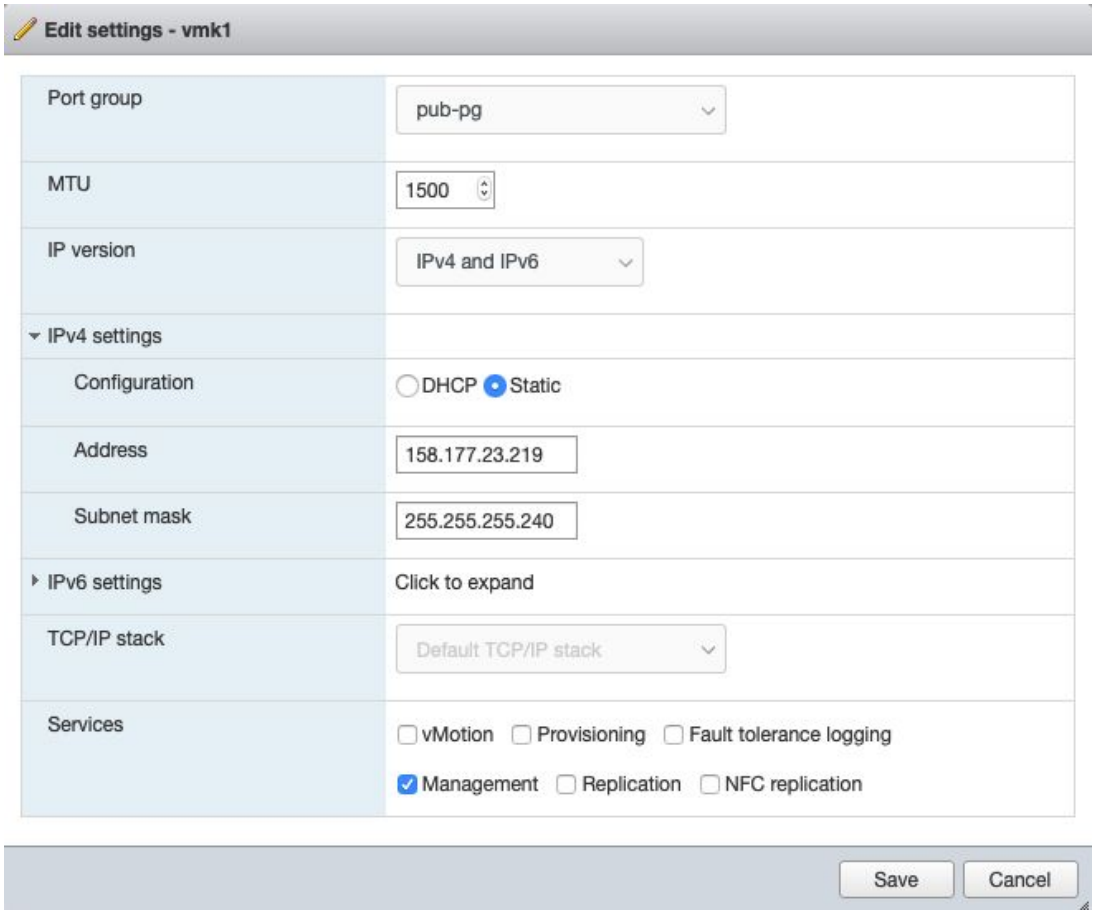

Run `esxcfg-route -a 10.0.0.0/8 [your servers private gateway ip]`

```
[root@baremetal01:~] esxcfg-route -a 10.0.0.0/8 10.194.251.65
```

Run `esxcfg-route [your servers public gateway]` to ensure that the public gateway is the default

```
[root@baremetal01:~] esxcfg-route 158.177.23.209
VMkernel default gateway set to 158.177.23.209
```

Troubleshooting:

```
esxcli network nic list
```

Make sure you have the right physical NICs , if your machine is not ordered dual, one public and one private are up.
So typically you should find 4 entries, 2 of them are up. Like below

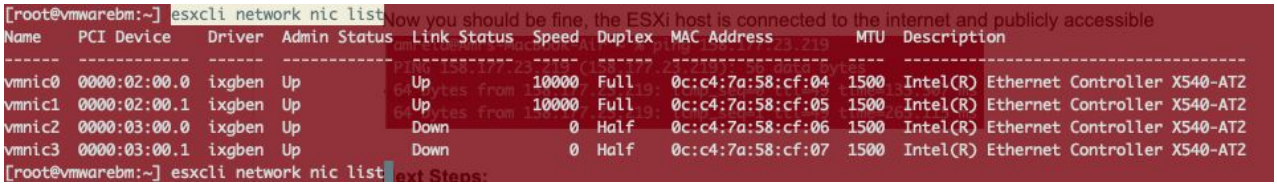

Now you got all set in the ESXi host, the ESXi host is connected to the internet and publicly accessible

```
amreid@Amrs-MacBook-Air ~ % ping 158.177.23.219
PING 158.177.23.219 (158.177.23.219): 56 data bytes
64 bytes from 158.177.23.219: icmp_seq=0 ttl=49 time=133.307 ms
64 bytes from 158.177.23.219: icmp_seq=1 ttl=49 time=265.115 ms
```

Provisioning the environment

The list appears in the screenshot below is what we need to achieve, I will explain how to create the first VM only.

Virtual machine ▲	Status	Used space	Guest OS
dns-server	Normal	104.11 GB	CentOS 7 (64-bit)
master01.eid.com	Normal	432.11 GB	CentOS 7 (64-bit)
worker-1.eid.com	Normal	216.11 GB	CentOS 7 (64-bit)
worker02.eid.com	Normal	116.11 GB	CentOS 7 (64-bit)

As a **prerequisite**,, we need to get a CentOS image to have it on the datastore.

```
cd vmfs/volumes/datastore1/iso/
wget
http://mirror.fra10.de.leaseweb.net/centos/7/isos/x86_64/CentOS-7-x86_64-Minimal-1908.iso
```

Secondly: Let's us create a VM

First we need to define if the VM will have private / public network adapters or both.

So, we need to add the related port-groups like below

Port groups | Virtual switches | Physical NICs | VMkernel NICs | TCP/IP stacks | Firewall rules

Add port group | Edit settings | Refresh | Actions

Name	Active ports	VLAN ID	Type	vSwitch
vm-private	0	0	Standard port group	vSwitch0
Management Network	0	0	Standard port group	vSwitch0
pg-private	1	0	Standard port group	vSwitch0
vm-pub	0	0	Standard port group	vSwitch1
pg-public	1	0	Standard port group	vSwitch1

Then , go ahead and create the VM using the wizard, default setting till you reach "Customize Settings", Add Network Adapter , In the network adapter pick the related port-group for each adapter.

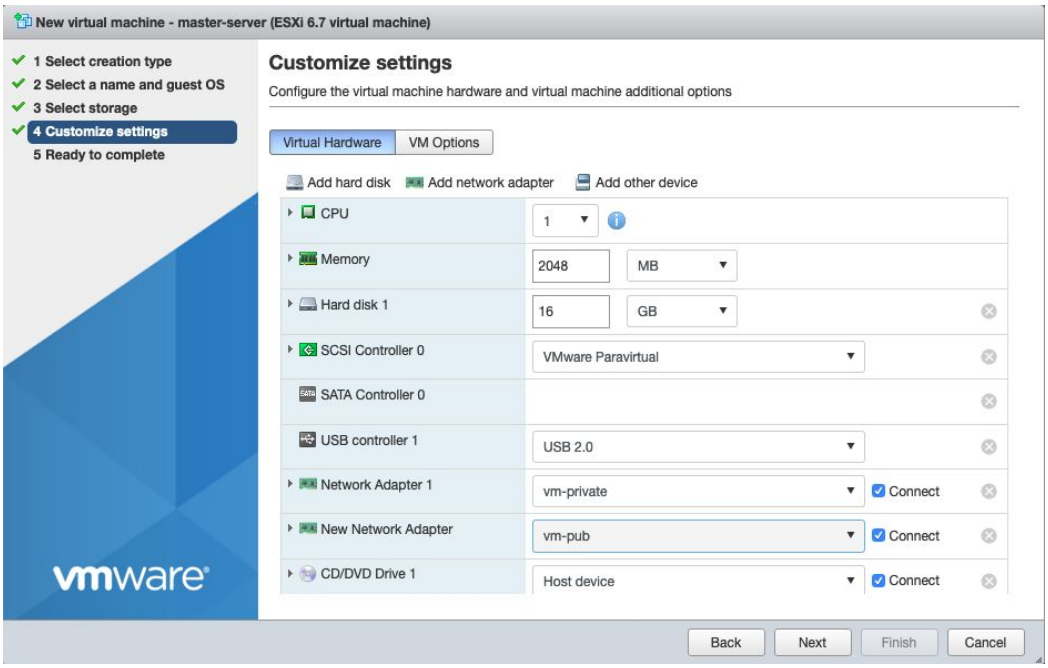

Following the same wizard, From the above step, From CD/DVD pick the proper iso for your OS. Complete the steps and power the VM on.

You can either configure the network/DNS before starting the real installation using CentOS setup wizard. For the Network, you will only be asked to provide an IP/Netmask and Gateway.
Or
After installing the OS, you can manually configure the network by accessing the VM from ESXi host, and using the CLi you can configure the network with a tool like nmtui

The below is to validate your network settings

- You can try to ping a domain like www.ibm.com
- If it worked, then you got all set. If not, look to DNS setup
- For now, you an use any public DNS like 8.8.8.8
- Try `cat /etc/resolv.conf` , if it is blank, you need to set up a DNS resolver. I would suggest editing the resolv.conf with `nameserver 8.8.8.8`
- Try ping `www.ibm.com` or `yum update` again.

One final step, as you will have to install many VMs, it is good to change the hostname for every VM to be descriptive, for CentOS, you can do it from ESXi installation wizard and you can do it later through

```
hostnamectl set-hostname "eidns.amreid.com"
exec bash
```

Section 03- Installing Openshift

03-01 Host Preparation - DNS Node Setup
DNSMASQ: How to install and configure

Install

```
yum install dnsmasq bind-utils
```

Configure

Edit those two files, and add the below line to both of them
/etc/resolv.conf and /etc/resolv.dnsmasq

```
nameserver 127.0.0.1
```

You will need to edit those two files 1- /etc/dnsmasq.conf and 2-/etc/hosts as follows

```
cp /etc/dnsmasq.conf /etc/dnsmasq.conf.orig
vi /etc/dnsmasq.conf
```

At the end of the file you can just add

```
server=8.8.8.8    #Google's nameserver
local=/amreid.com/    #you domain servers to be looked locally
#Log-queries         # In case you need to enable logs
#log-facility=/var/log/dnsmasq.log
```

Now, edit vi /etc/hosts and add your hosts like below

Then, start and enable the dns service.

```
systemctl start dnsmasq
systemctl enable dnsmasq
systemctl status dnsmasq
```

You should expect results like below saying that it is up and running.

Firewall

For now only, to keep the installation easier, let's either stop the firewall if you don't need it, Or open DNS and DHCP services in the firewall configuration, to allow requests from hosts on your LAN to pass to the dnsmasq server.

```
[root@dns-server ~]# firewall-cmd --add-service=dns --permanent
success
[root@dns-server ~]# firewall-cmd --add-service=dhcp --permanent
success
[root@dns-server ~]# firewall-cmd --reload
success
[root@dns-server ~]#
```

If you like to disable it you can use the below commands

```
systemctl status firewalld
service firewalld stop
systemctl disable firewalld
```

Validation

From worker or Master node, you can validate also by testing if you can reach to each other like

```
[root@master ~]# ping worker01.amreid.com
PING worker01.amreid.com (159.122.69.83) 56(84) bytes of data.
64 bytes from worker01.amreid.com (159.122.69.83): icmp_seq=1 ttl=64 time=0.135 ms
64 bytes from worker01.amreid.com (159.122.69.83): icmp_seq=2 ttl=64 time=0.120 ms
^C
--- worker01.amreid.com ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 999ms
rtt min/avg/max/mdev = 0.120/0.127/0.135/0.013 ms
[root@master ~]#
```

From the DNS server , from each other host, you can check if it can access the others.

```
dig worker01.amreid.com
OR  # nslookup worker01.amreid.com
```

Preferably nslookup and make sure to read the results carefully, If it shows this last line message, then it is not going to work properly.

```
[root@eidnsmasq ~]# nslookup master.amreid.com
Server:         127.0.0.1
Address:        127.0.0.1#53

Name:   master.amreid.com
Address: 159.122.69.84
** server can't find master.amreid.com: SERVFAIL
```

The right execution should looks like

```
[root@eidnsmasq ~]# nslookup worker02.amreid.com
Server:         127.0.0.1
Address:        127.0.0.1#53

Name:   worker02.amreid.com
Address: 159.122.69.86

[root@eidnsmasq ~]#
```

03-02 Installing Openshift -3.11

Host Preparation: Stage 1
Make sure that all the prerequisites are working fine. DNS is up and running , Masters and Worker nodes can see each other on top of the DNS.

Actions:
ssh-keygen on master node
distribute the key all over the nodes. (for host in xxxxxxx; do ssh-copy-id $host done)
On the Master Node:

```
ssh-keygen
```

Then

```
for host in master.amreid.com worker01.amreid.com worker02.amreid.com; do
ssh-copy-id $host; done
```

As a sanity check, you can validate the above step has been successfully by trying to ssh to each one of the worker nodes while being on the master node, you should login without any need to provide a password or a key

Host Preparation: Stage 2
Actions: Install base packages
On the master node: Install all the below packages

```
yum install wget git net-tools bind-utils yum-utils iptables-services
bridge-utils bash-completion kexec-tools sos psacct
yum install epel-release
yum install ansible pyOpenSSL
cd ~
git clone https://github.com/openshift/openshift-ansible
cd openshift-ansible
git checkout release-3.11
```

Host Preparation: Stage 3
Actions: Install docker
On the master node: run the below commands

```
yum install docker-1.13.1
```

This exact version is advised to be used with openshift3.11.
Here we need to add a dedicated LVM Volume Group , so if not added already, we can add a separate disk to the master VM then

`Lsblk` or `fdisk -l`

```
[root@master ~]# lsblk
NAME            MAJ:MIN RM   SIZE RO TYPE MOUNTPOINT
sda               8:0    0   500G  0 disk
|-sda1            8:1    0     1G  0 part /boot
`-sda2            8:2    0   499G  0 part
  |-centos-root 253:0    0    50G  0 lvm  /
  |-centos-swap 253:1    0   7.9G  0 lvm  [SWAP]
  `-centos-home 253:2    0 441.1G  0 lvm  /home
sdb               8:16   0   200G  0 disk
sr0              11:0    1  1024M  0 rom
```

```
[root@master ~]# fdisk -l | grep /dev
Disk /dev/sda: 536.9 GB, 536870912000 bytes, 1048576000 sectors
/dev/sda1   *        2048     2099199     1048576   83  Linux
/dev/sda2         2099200  1048575999   523238400   8e  Linux LVM
Disk /dev/mapper/centos-root: 53.7 GB, 53687091200 bytes, 104857600 sectors
Disk /dev/mapper/centos-swap: 8455 MB, 8455716864 bytes, 16515072 sectors
Disk /dev/mapper/centos-home: 473.6 GB, 473645973504 bytes, 925089792 sectors
Disk /dev/sdb: 214.7 GB, 214748364800 bytes, 419430400 sectors
/dev/sdb1            2048   419430399   209714176   8e  Linux LVM
```

To get the name of the disk, here it is /dev/sdb

- Edit the file `vi /etc/sysconfig/docker-storage-setup`
  Add
  `STORAGE_DRIVER=overlay2`
  `DEVS=/dev/sdb`
  `VG=docker-vg`
  Run the file using the below command
- `docker-storage-setup create dockervg /etc/sysconfig/docker-storage-setup`

```
[root@master ~]# docker-storage-setup create dockervg /etc/sysconfig/docker-storage-setup
INFO: Writing zeros to first 4MB of device /dev/sdb
4+0 records in
4+0 records out
4194304 bytes (4.2 MB) copied, 0.510424 s, 8.2 MB/s
INFO: Device node /dev/sdb1 exists.
  Physical volume "/dev/sdb1" successfully created.
  Volume group "docker-vg" successfully created
```

In order to verify this phase, just run `vgs` to see the created Volume Groups. (20G is enough for the dedicated VG)

```
[root@master ~]# vgs
  VG        #PV #LV #SN Attr   VSize    VFree
  centos      1   3   0 wz--n- <499.00g    4.00m
  docker-vg   1   0   0 wz--n- <200.00g <200.00g
```

Openshift Installation: Using Ansible Installer

Actions:

- Installing Openshift using Ansible Playbook
- The initial configuration is in Ansible Inventory file /etc/ansible/hosts (on the Master Node).

Step 01

```
cd ~/openshift-ansible
cp hosts.localhost /etc/ansible/hosts
vi /etc/ansible/hosts
```

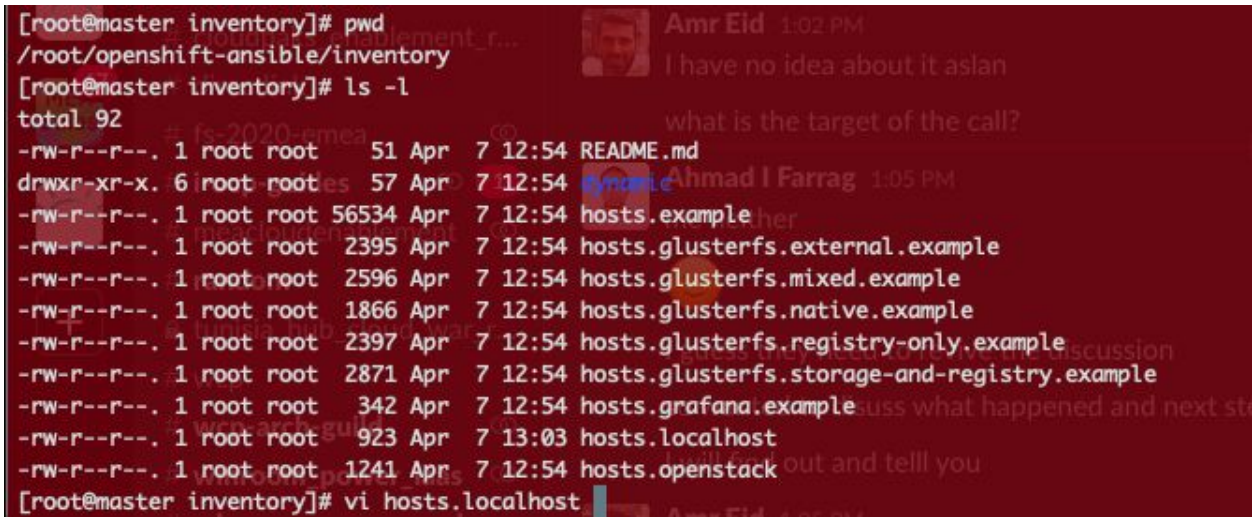

Step 02

Edit your ansible hosts file to have your nodes under the proper sections and make sure that you've added group names like below.

```
openshift_deployment_type=origin
openshift_portal_net=172.30.0.0/16
openshift_disable_check=disk_availability,memory_availability

[masters]
master01.eid.com

[etcd]
master01.eid.com

[nodes]
master01.eid.com openshift_node_group_name='node-config-master'
worker01.eid.com openshift_node_group_name='node-config-compute'
infra01.eid.com openshift_node_group_name='node-config-infra'
```

```
#bare minimum hostfile

[OSEv3:children]
masters
nodes
etcd

[OSEv3:vars]
# if your target hosts are Fedora uncomment this
#ansible_python_interpreter=/usr/bin/python3
openshift_deployment_type=origin
openshift_portal_net=172.30.0.0/16
# localhost likely doesn't meet the minimum requirements
openshift_disable_check=disk_availability,memory_availability

#openshift_node_groups=[{'name': 'node-config-all-in-one', 'labels': ['node-role.kubernetes.io/master=true', 'nod
ue']}]


[masters]
master01.eid.com
#localhost ansible_connection=local

[etcd]
master01.eid.com
#localhost ansible_connection=local

[nodes]
master01.eid.com openshift_node_group_name='node-config-master'
worker01.eid.com openshift_node_group_name='node-config-compute'
infra01.eid.com openshift_node_group_name='node-config-infra'

# openshift_node_group_name should refer to a dictionary with matching key of name in list openshift_node_groups.
#localhost ansible_connection=local openshift_node_group_name="node-config-all-in-one"
```

Step 03

Make sure that you are on the openshift-ansible root folder then run the prerequisites.yml script

```
ansible-playbook playbooks/prerequisites.yml
```

This prerequisites script is testing the environment before the installation. You should receive green results like below.

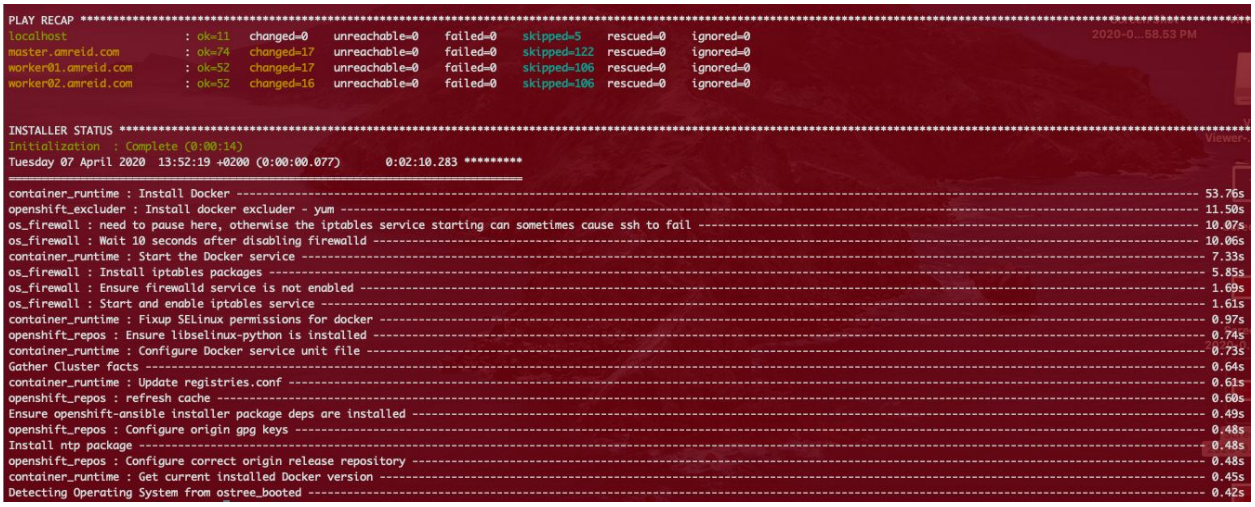

```
PLAY RECAP *********************************************************************
localhost                  : ok=11   changed=0    unreachable=0    failed=0    skipped=5    rescued=0    ignored=0
master.amreid.com          : ok=74   changed=17   unreachable=0    failed=0    skipped=122  rescued=0    ignored=0
worker01.amreid.com        : ok=52   changed=17   unreachable=0    failed=0    skipped=106  rescued=0    ignored=0
worker02.amreid.com        : ok=52   changed=16   unreachable=0    failed=0    skipped=106  rescued=0    ignored=0


INSTALLER STATUS ***************************************************************
Initialization  : Complete (0:00:14)
Tuesday 07 April 2020  13:52:19 +0200 (0:00:00.077)       0:02:10.283 *********
===============================================================================
container_runtime : Install Docker ---------------------------------------- 53.76s
openshift_excluder : Install docker excluder - yum ------------------------ 11.50s
os_firewall : need to pause here, otherwise the iptables service starting can sometimes cause ssh to fail ---------- 10.07s
os_firewall : Wait 10 seconds after disabling firewalld ------------------- 10.06s
container_runtime : Start the Docker service ------------------------------- 7.33s
os_firewall : Install iptables packages ------------------------------------ 5.85s
os_firewall : Ensure firewalld service is not enabled ---------------------- 1.69s
os_firewall : Start and enable iptables service ---------------------------- 1.61s
container_runtime : Fixup SELinux permissions for docker ------------------- 0.97s
openshift_repos : Ensure libselinux-python is installed -------------------- 0.74s
container_runtime : Configure Docker service unit file --------------------- 0.73s
Gather Cluster facts ------------------------------------------------------- 0.64s
container_runtime : Update registries.conf --------------------------------- 0.61s
openshift_repos : refresh cache -------------------------------------------- 0.60s
Ensure openshift-ansible installer package deps are installed -------------- 0.49s
openshift_repos : Configure origin gpg keys -------------------------------- 0.48s
Install ntp package -------------------------------------------------------- 0.48s
openshift_repos : Configure correct origin release repository -------------- 0.48s
container_runtime : Get current installed Docker version ------------------- 0.45s
Detecting Operating System from ostree_booted ------------------------------ 0.42s
```

Step 04

```
ansible-playbook playbooks/deploy_cluster.ym
```

This is going to take time (~15min), then you should receive green results similar to the previous step.

```
PLAY RECAP *********************************************************************
infra01.eid.com            : ok=115  changed=67   unreachable=0    failed=0
localhost                  : ok=11   changed=0    unreachable=0    failed=0
master01.eid.com           : ok=564  changed=148  unreachable=0    failed=0
worker01.eid.com           : ok=108  changed=20   unreachable=0    failed=0


INSTALLER STATUS ***************************************************************
Initialization               : Complete (0:00:15)
Health Check                 : Complete (0:00:20)
Node Bootstrap Preparation   : Complete (0:04:46)
etcd Install                 : Complete (0:00:23)
Master Install               : Complete (0:01:48)
Master Additional Install    : Complete (0:00:24)
Node Join                    : Complete (0:00:31)
Hosted Install               : Complete (0:00:29)
Cluster Monitoring Operator  : Complete (0:01:12)
Web Console Install          : Complete (0:01:25)
Console Install              : Complete (0:01:16)
Service Catalog Install      : Complete (0:03:59)
Friday 10 April 2020  16:27:38 +0200 (0:00:00.032)       0:17:11.795 **********
```

Step 05 - Validating and Testing your open shift environment

With the previous good news, you can add your master node IP to your hosts file and then using https open your browser to access openshift console, note that openshift 3.11 is using port 8433.

https://master01.eid.com:8443

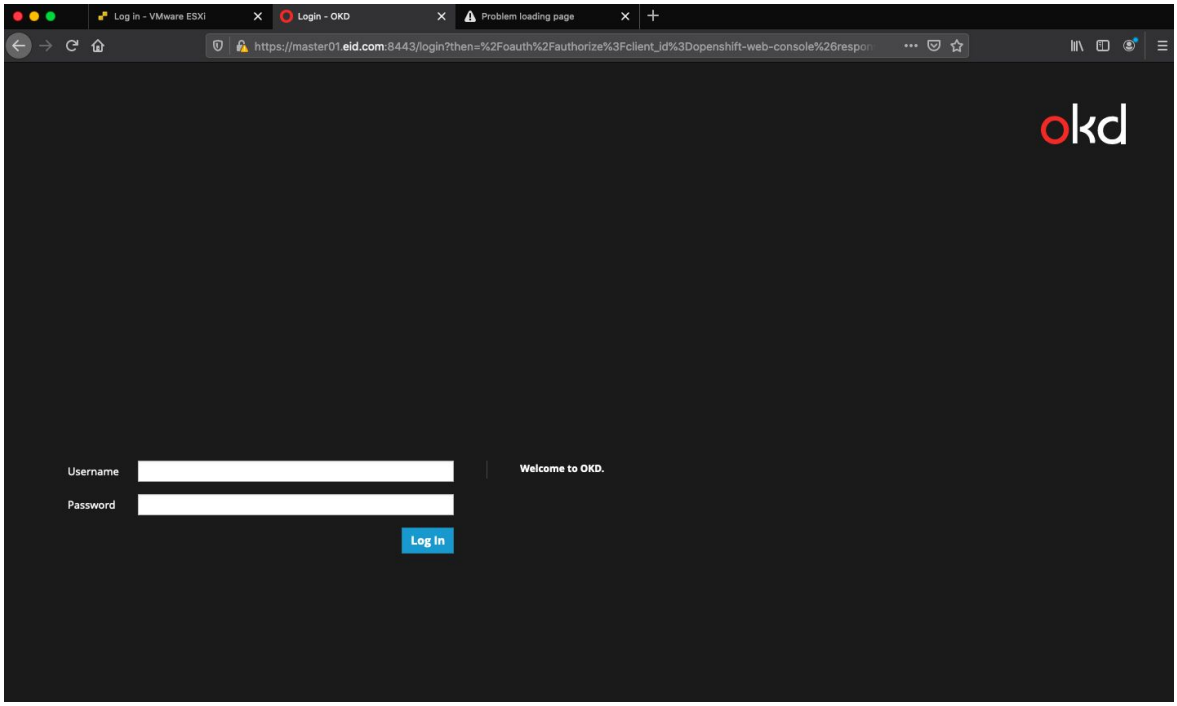

You can try log in using developer / any password

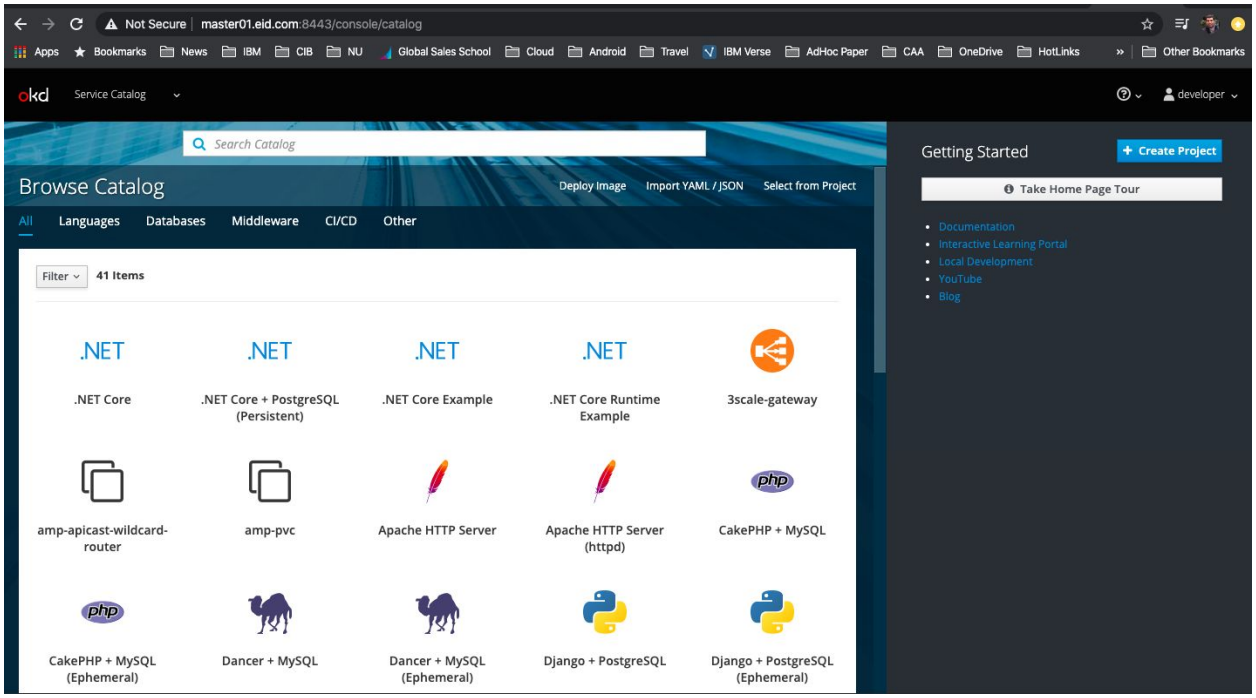

Troubleshooting

This section is trying to capture some of the errors which may happen while trying to install specially for your first time.

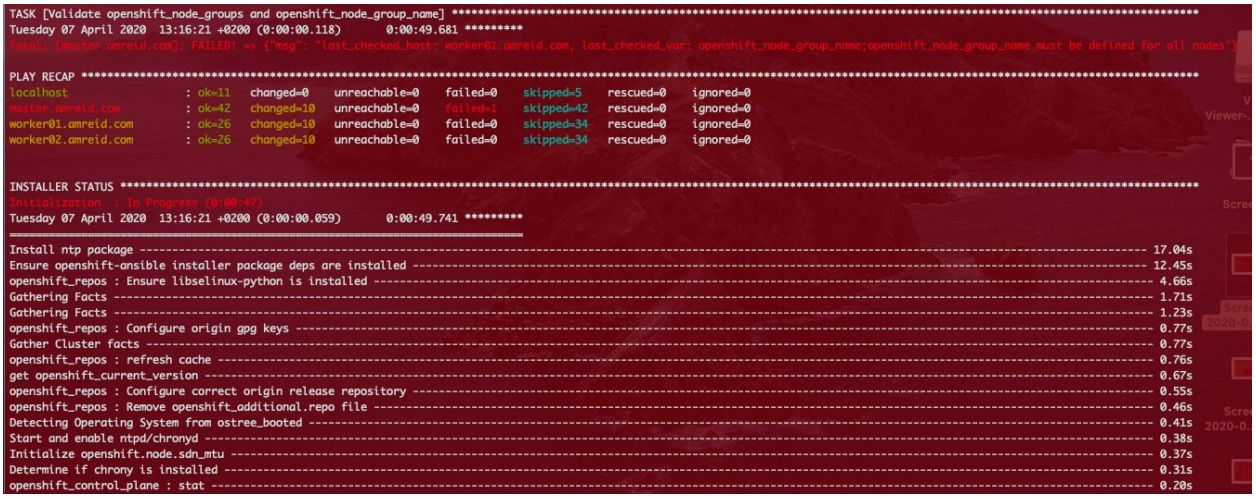

```
TASK [openshift_control_plane : Report control plane errors] ********************************************************************
Tuesday 07 April 2020  14:30:03 +0200 (0:00:00.097)       0:30:50.099 *********
fatal: [master.amreid.com]: FAILED! => {"changed": false, "msg": "Control plane pods didn't come up and become ready"}

NO MORE HOSTS LEFT **************************************************************************************************************

PLAY RECAP **********************************************************************************************************************
localhost                  : ok=11   changed=0    unreachable=0    failed=0    skipped=5    rescued=0    ignored=0
master.amreid.com          : ok=305  changed=155  unreachable=0    failed=1    skipped=232  rescued=0    ignored=4
worker01.amreid.com        : ok=106  changed=64   unreachable=0    failed=0    skipped=94   rescued=0    ignored=0
worker02.amreid.com        : ok=106  changed=64   unreachable=0    failed=0    skipped=94   rescued=0    ignored=0


INSTALLER STATUS ****************************************************************************************************************
Initialization             : Complete (0:00:10)
Health Check               : Complete (0:00:30)
Node Bootstrap Preparation : Complete (0:09:10)
etcd Install               : Complete (0:00:25)
Master Install             : In Progress (0:20:34)
        This phase can be restarted by running: playbooks/openshift-master/config.yml
Tuesday 07 April 2020  14:30:03 +0200 (0:00:00.042)       0:30:50.142 *********
```

In order to fix these kinds of errors, Here are some notes you can keep in mind.

- Try using `-vvv` to run the scripts in verbose mode if you want to have more descriptive errors.
- Any connection related errors (e.x SSH connection to any node), most probably it will be related to the host preparation section in this document so, give Dnsmasq setup your main focus. I've listed it above in detail.
- A lot of ansible script issues will be related to missing some detail in Ansible hosts file. So review the hosts file carefully and make sure that the group name is setted for all the nodes.
- Before initiating the deploy script, make sure you can access all the nodes from Master Node using Ansible like below.
- `ansible -m ping all`
- The results should be looks like

```
[root@master01 openshift-ansible]# ansible -m ping all
master01.eid.com | SUCCESS => {
    "changed": false,
    "ping": "pong"
}
worker01.eid.com | SUCCESS => {
    "ansible_facts": {
        "discovered_interpreter_python": "/usr/bin/python"
    },
    "changed": false,
    "ping": "pong"
}
[root@master01 openshift-ansible]#
```

- You may need to install Ansible stable version well tested with Openshift 3.1 recommended version is: 2.7.1
  How to install Ansible 2.7.1

```
curl https://bootstrap.pypa.io/get-pip.py -o get-pip.py
python get-pip.py
pip install ansible==2.7.10
ansible --version
```

```
[root@master01 openshift-ansible]# ansible --version
ansible 2.7.10
  config file = /root/openshift-ansible/ansible.cfg
  configured module search path = [u'/root/.ansible/plugins/modules', u'/usr/share/ansible/plugins/modules']
  ansible python module location = /usr/lib/python2.7/site-packages/ansible
  executable location = /usr/bin/ansible
  python version = 2.7.5 (default, Aug  7 2019, 00:51:29) [GCC 4.8.5 20150623 (Red Hat 4.8.5-39)]
[root@master01 openshift-ansible]#
```

- You may face some issues while logging into the dns server like below

```
@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@

@    WARNING: REMOTE HOST IDENTIFICATION HAS CHANGED!     @

@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@@
```

How to fix it, simply run the below command.

```
ssh-keygen -R
```